

## Homework 03: You will run Logistic Regression, K-Nearest

- ✓ Neighbor, Decision Tree, and Random Forest Classifier to predict survival for the Titanic Dataset.

Then, you will check and print the performance of your model.

---

- ✓ First, get all your required packages. Note: the list below is not exhaustive, if you need more packages, please import them as needed.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import sklearn.metrics as metrics
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```


- ✓ We are going to obtain the Titanic dataset from DataScienceDojo's github page. Thank you to them for the readily available data.

Here is the link: <https://github.com/datasciencedojo/datasets/blob/master/titanic.csv>

Import the file as a DataFrame called `titanic`.

```
# https://github.com/datasciencedojo/datasets/blob/master/titanic.csv

titanic = pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv')
titanic.head()
```



	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38.0	1	0	PC 17599

✓ **Survived** is your target variable, also known as your dependent variable.

Your attributes/features/independent variables will help predict survival on the Titanic.

But first, you need to preprocess the data.


- Note: **Survived** is your target variable (**Y**).
- **Pclass**, **Sex**, **Age**, **SibSp**, **Parch**, and **Fare** will certainly be important predictors of whether a passenger survived or not. Hence, they will be included in your attributes list (**X**).
- The **Name**, **Ticket**, and **Cabin** are not useful features. Someone's name has no bearing on whether they survive or not. Similarly, a ticket number is just a unique identifier for a passenger - it is not meaningful, ordered data. So we can drop these 3. **NOTE**: NEVER drop variables from the original dataset. Either create a new df for relevant features, or create a copy of titanic and then drop the ones you do not want.
- Where they **Embarked** can be meaningful, but the data is a string variable. Let us convert it to an integer. This can be done with **np.where** or with **label encoding**. I will help you with this step. You have to do the rest of the preprocessing steps.

```

titanic.Embarked.unique()
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
titanic['embarked'] = le.fit_transform(titanic['Embarked'])

titanic.head()

```



	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th	female	38.0	1	0	PC 17599 7

✓ Perfect! Now You can use Pclass, Sex, Age, SibSp, Parch, Fare, and embarked as your features.

Step 1: First shuffle your dataset. Step 2: Create X and Y arrays. You can refer to the class file for this. Y is the target (single column), X comprises all the relevant features.

```
titanic = titanic.sample(frac=1, random_state=42).reset_index(drop=True)
```

```

X = titanic[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'embarked']]
y = titanic['Survived']

```

## ✓ 1. Logistic Regression

Use this model to predict survival on the Titanic.

### Part I:

First, use a 80:20 train-test split. Run your logistic regression prediction model.

Then, report the accuracy, precision, recall, f-score, sensitivity, specificity, and the confusion matrix. Plot the ROC curve.

Part II: Repeat all the above with a 60:40 split.

Compare the results between Part 1 and Part 2. Which split gave you better results in your opinion?

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.impute import SimpleImputer

# X = titanic[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'embarked']]
y = titanic['Survived']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s

# Convert 'Sex' column to numerical values *before* imputation
X_train['Sex'] = np.where(X_train['Sex'] == 'female', 1, 0)
X_test['Sex'] = np.where(X_test['Sex'] == 'female', 1, 0)

# Impute missing values using the mean
imputer = SimpleImputer(strategy='mean') # or strategy='median', etc.
X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)

# Convert back to DataFrame for easier handling if needed
X_train = pd.DataFrame(X_train, columns=['Pclass', 'Sex', 'Age', 'SibSp', 'Parch']
X_test = pd.DataFrame(X_test, columns=['Pclass', 'Sex', 'Age', 'SibSp', 'Parch',

# Now you can proceed with fitting the model
model = LogisticRegression(max_iter=1000) # Increase max_iter if needed
model.fit(X_train, y_train)

```



▼ LogisticRegression ⓘ ?  
 LogisticRegression(max\_iter=1000)

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_sco

```

```

# Assuming X and y are already defined as in your previous code

# --- 80:20 Split ---
X_train_80, X_test_80, y_train_80, y_test_80 = train_test_split(X, y, test_size=0.2)
X_train_80['Sex'] = np.where(X_train_80['Sex'] == 'female', 1, 0)
X_test_80['Sex'] = np.where(X_test_80['Sex'] == 'female', 1, 0)
imputer_80 = SimpleImputer(strategy='mean')
X_train_80 = imputer_80.fit_transform(X_train_80)
X_test_80 = imputer_80.transform(X_test_80)
model_80 = LogisticRegression(max_iter=1000)
model_80.fit(X_train_80, y_train_80)
y_pred_80 = model_80.predict(X_test_80)

# --- 60:40 Split ---
X_train_60, X_test_60, y_train_60, y_test_60 = train_test_split(X, y, test_size=0.4)
X_train_60['Sex'] = np.where(X_train_60['Sex'] == 'female', 1, 0)
X_test_60['Sex'] = np.where(X_test_60['Sex'] == 'female', 1, 0)
imputer_60 = SimpleImputer(strategy='mean')
X_train_60 = imputer_60.fit_transform(X_train_60)
X_test_60 = imputer_60.transform(X_test_60)
model_60 = LogisticRegression(max_iter=1000)
model_60.fit(X_train_60, y_train_60)
y_pred_60 = model_60.predict(X_test_60)

# --- Evaluation and Comparison ---
def evaluate_model(y_true, y_pred, split_name):
    print(f"--- {split_name} Split ---")
    print("Accuracy:", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred))
    print("Recall:", recall_score(y_true, y_pred))
    print("F1-score:", f1_score(y_true, y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_true, y_pred))
    print("Classification Report:\n", classification_report(y_true, y_pred))
    print("\n")

evaluate_model(y_test_80, y_pred_80, "80:20")
evaluate_model(y_test_60, y_pred_60, "60:40")

# --- Conclusion ---
# Compare the evaluation metrics (accuracy, precision, recall, F1-score)
# for both splits and decide which one performed better.
# Consider the trade-off between having more data for training (80:20)
# and having a larger test set for more robust evaluation (60:40).

```



--- 80:20 Split ---

Accuracy: 0.7430167597765364  
 Precision: 0.6428571428571429  
 Recall: 0.6818181818181818  
 F1-score: 0.6617647058823529  
 Confusion Matrix:

```
[[88 25]
 [21 45]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.78	0.79	113
1	0.64	0.68	0.66	66
accuracy			0.74	179
macro avg	0.73	0.73	0.73	179
weighted avg	0.75	0.74	0.74	179

--- 60:40 Split ---

Accuracy: 0.7759103641456583  
 Precision: 0.6845637583892618  
 Recall: 0.7555555555555555  
 F1-score: 0.7183098591549296  
 Confusion Matrix:

```
[[175 47]
 [ 33 102]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.79	0.81	222
1	0.68	0.76	0.72	135
accuracy			0.78	357
macro avg	0.76	0.77	0.77	357
weighted avg	0.78	0.78	0.78	357



## ✓ Now K Nearest Neighbors:

### 2. KNN: Use this model to predict survival on the Titanic.

Part I:

Use a 80:20 train-test split. Run your KNN choosing 3 nearest neighbors.

Then, report the accuracy, precision, recall, f-score, sensitivity, specificity, and the confusion matrix. Plot the ROC curve.

Part II: Repeat the above with neighbors = 5.

Compare the results between Part 1 and Part 2. Which neighbor selection gave you better results?

# start working here. Feel free to use several separate blocks of code.

```
# convert Sex to 1 and 0 where 1 is female
titanic['Sex'] = np.where(titanic['Sex'] == 'female', 1, 0)
X = titanic[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'embarked']]
y = titanic['Survived']
```

## ✓ Finally, Tree-Methods:

### 3. Decision Tree: Use this model to predict survival on the Titanic.

Part I:

Use a 80:20 train-test split.

Then, report the accuracy, precision, recall, f-score, sensitivity, specificity, and the confusion matrix. **Plot** the ROC curve.

Part II: Repeat the above with a 50:50 train test split.

Compare the results between Part 1 and Part 2. Which split gave you better results?

Try a Random Forest Classifier as well. Works very similarly to how a decision tree does.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt

# Assuming X and y are already defined

# --- Decision Tree ---

# 80:20 Split
X_train_80, X_test_80, y_train_80, y_test_80 = train_test_split(X, y, test_size=0.2)
dt_80 = DecisionTreeClassifier(random_state=42)
dt_80.fit(X_train_80, y_train_80)
y_pred_dt_80 = dt_80.predict(X_test_80)

# 50:50 Split
X_train_50, X_test_50, y_train_50, y_test_50 = train_test_split(X, y, test_size=0.5)
dt_50 = DecisionTreeClassifier(random_state=42)
dt_50.fit(X_train_50, y_train_50)
y_pred_dt_50 = dt_50.predict(X_test_50)

# --- Random Forest ---

# 80:20 Split
rf_80 = RandomForestClassifier(random_state=42)
rf_80.fit(X_train_80, y_train_80)
y_pred_rf_80 = rf_80.predict(X_test_80)

# 50:50 Split
rf_50 = RandomForestClassifier(random_state=42)
rf_50.fit(X_train_50, y_train_50)
y_pred_rf_50 = rf_50.predict(X_test_50)

# --- Evaluation ---

def evaluate_model(y_true, y_pred, split_name, model_name):
    print(f"--- {model_name} - {split_name} Split ---")
    print("Accuracy:", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred))
    print("Recall:", recall_score(y_true, y_pred))
    print("F1-score:", f1_score(y_true, y_pred))
```

```

# ... (add sensitivity, specificity, confusion matrix calculations here) ...
print("Confusion Matrix:\n", confusion_matrix(y_true, y_pred))
print("Classification Report:\n", classification_report(y_true, y_pred))
print("\n")

# ROC Curve
fpr, tpr, thresholds = roc_curve(y_true, y_pred)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)'
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve - {model_name} - {split_name} Split')
plt.legend(loc="lower right")
plt.show()

evaluate_model(y_test_80, y_pred_dt_80, "80:20", "Decision Tree")
evaluate_model(y_test_50, y_pred_dt_50, "50:50", "Decision Tree")
evaluate_model(y_test_80, y_pred_rf_80, "80:20", "Random Forest")
evaluate_model(y_test_50, y_pred_rf_50, "50:50", "Random Forest")

# --- Conclusion ---
# Compare the evaluation metrics and ROC curves for both splits
# and both models to determine which combination performed better.

```



--- Decision Tree - 80:20 Split ---

Accuracy: 0.7653631284916201

Precision: 0.6818181818181818

Recall: 0.6818181818181818

F1-score: 0.6818181818181818


Confusion Matrix:

```
[[92 21]
```

```
[21 45]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.81	0.81	113
1	0.68	0.68	0.68	66
accuracy			0.77	179
macro avg	0.75	0.75	0.75	179
weighted avg	0.77	0.77	0.77	179



--- Decision Tree - 50:50 Split ---

Accuracy: 0.7488789237668162

Precision: 0.6612903225806451

Recall: 0.7151162790697675

F1-score: 0.6871508379888268

Confusion Matrix:

```
[[211  63]
```

```
[ 49 123]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.77	0.79	274
1	0.66	0.72	0.69	172
accuracy			0.75	446
macro avg	0.74	0.74	0.74	446
weighted avg	0.75	0.75	0.75	446

### ROC Curve - Decision Tree - 50:50 Split

